

```

In [30... import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# EDA
data = pd.read_csv('/content/Brain Tumor.csv')

print(data.describe())

sns.countplot(x='Class', data=data)
plt.title('Distribution of Classes')
plt.show()

# corr m - remind yourself to remove this if it does not provide useful in
corr_matrix = data.corr()
plt.figure(figsize=(10, 10))
sns.heatmap(corr_matrix, annot=True, fmt=".2f", cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()


selected_features = ["Mean", "Variance", "Standard Deviation", "Entropy"]
sns.pairplot(data[selected_features])
plt.show()

```

	Class	Mean	Variance	Standard Deviation	Entropy \
count	3762.000000	3762.000000	3762.000000	3762.000000	3762.000000
mean	0.447368	9.488890	711.101063	25.182271	0.073603
std	0.497288	5.728022	467.466896	8.773526	0.070269
min	0.000000	0.078659	3.145628	1.773592	0.000882
25%	0.000000	4.982395	363.225459	19.058475	0.006856
50%	0.000000	8.477531	622.580417	24.951560	0.066628
75%	1.000000	13.212723	966.954319	31.095889	0.113284
max	1.000000	33.239975	2910.581879	53.949809	0.394539


	Skewness	Kurtosis	Contrast	Energy	ASM \
count	3762.000000	3762.000000	3762.000000	3762.000000	3762.000000
mean	4.102727	24.389071	127.961459	0.204705	0.058632
std	2.560940	56.434747	109.499601	0.129352	0.058300
min	1.886014	3.942402	3.194733	0.024731	0.000612
25%	2.620203	7.252852	72.125208	0.069617	0.004847
50%	3.422210	12.359088	106.737418	0.225496	0.050849
75%	4.651737	22.640304	161.059006	0.298901	0.089342
max	36.931294	1371.640060	3382.574163	0.589682	0.347725

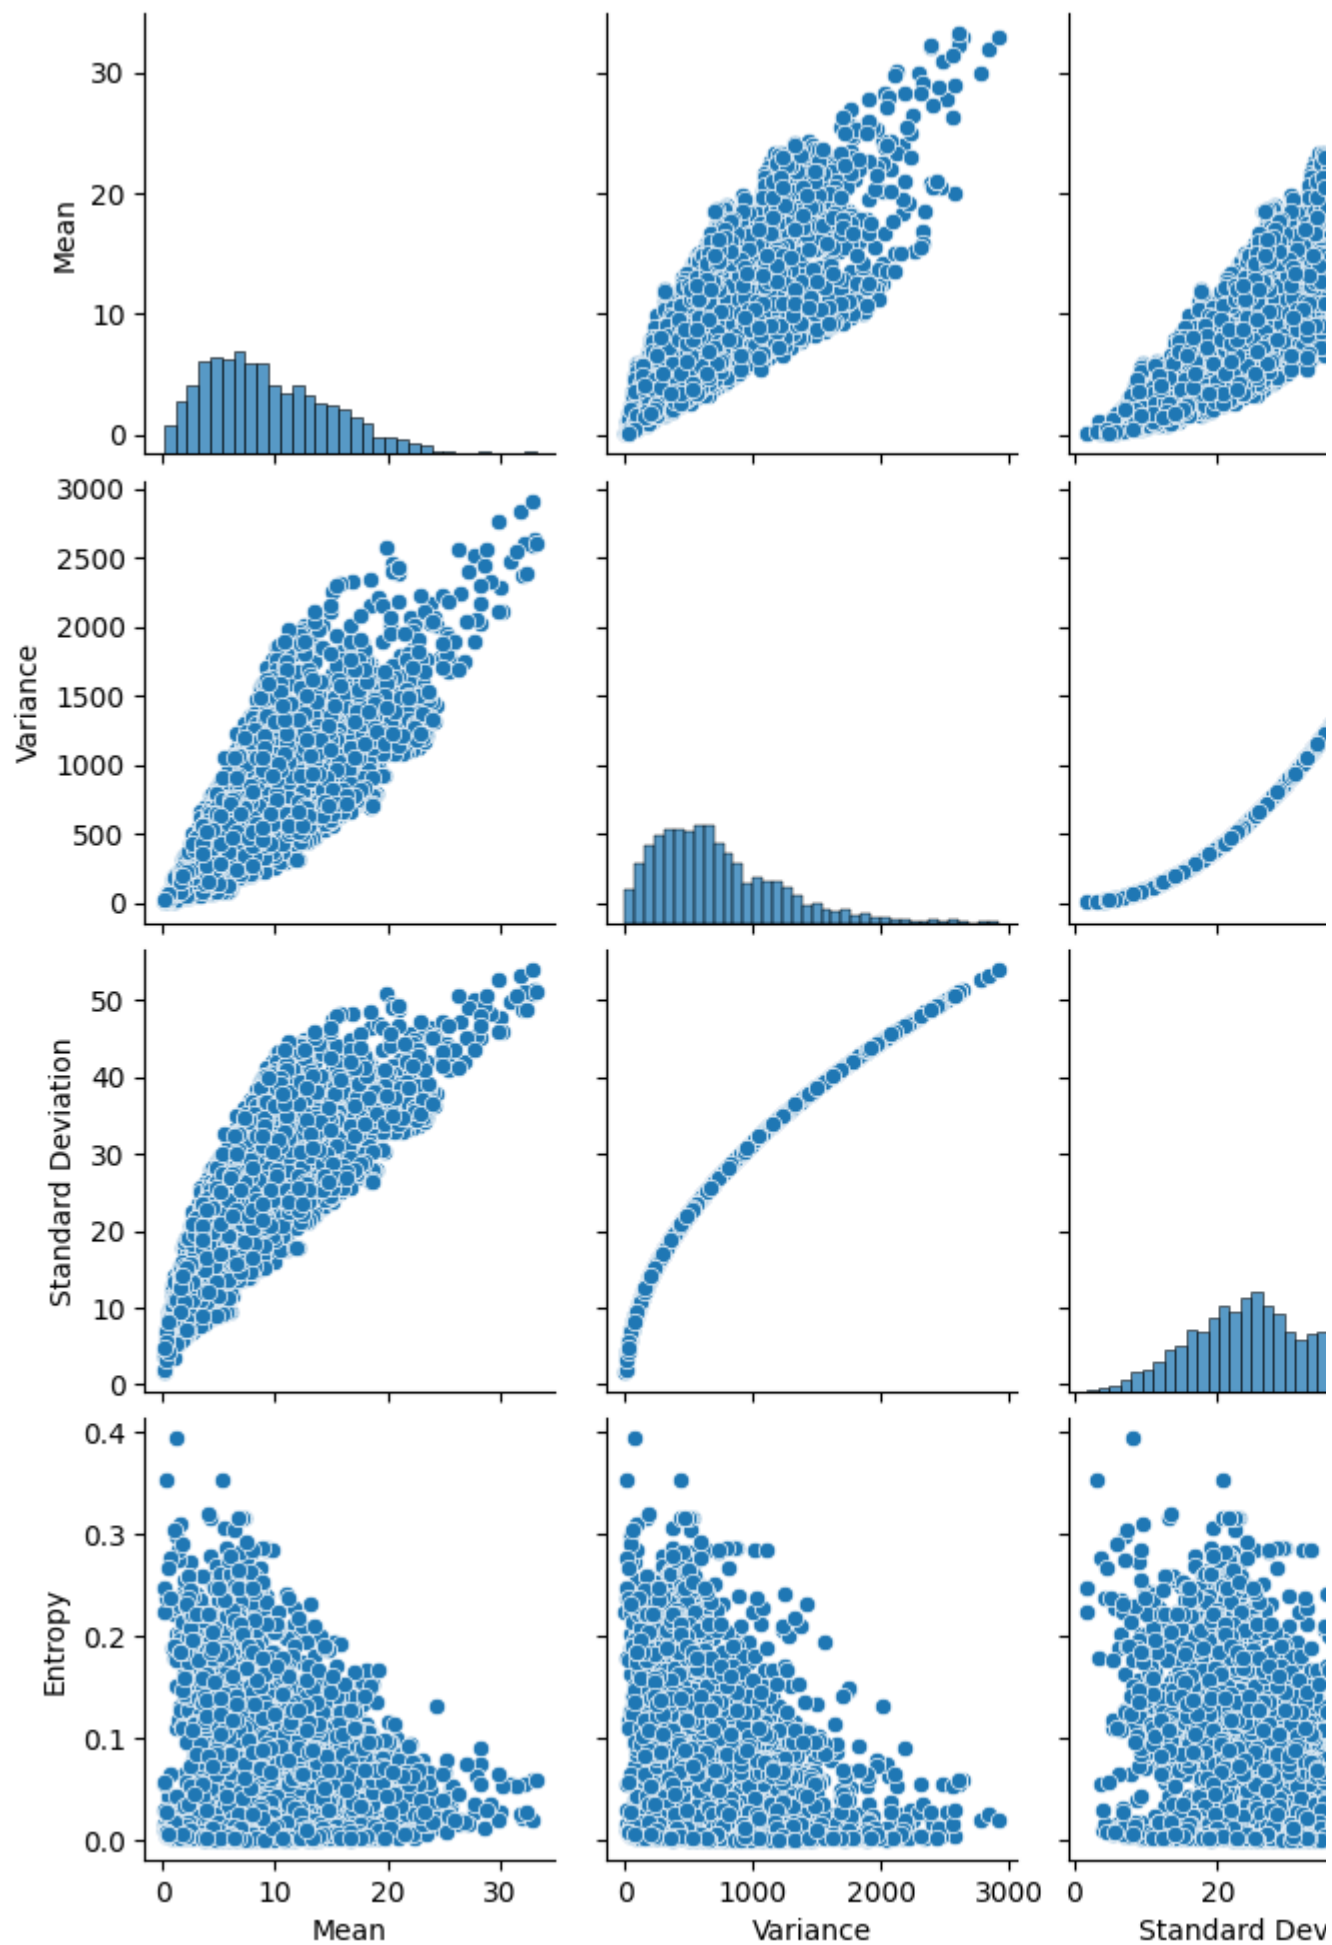
	Homogeneity	Dissimilarity	Correlation	Coarseness
count	3762.000000	3762.000000	3762.000000	3.762000e+03
mean	0.479252	4.698498	0.955767	7.458341e-155
std	0.127929	1.850173	0.026157	0.000000e+00
min	0.105490	0.681121	0.549426	7.458341e-155
25%	0.364973	3.412363	0.947138	7.458341e-155
50%	0.512551	4.482404	0.961610	7.458341e-155
75%	0.575557	5.723821	0.971355	7.458341e-155
max	0.810921	27.827751	0.989972	7.458341e-155



```
<ipython-input-30-870efe3d9939>:17: FutureWarning: The default value of  
numeric_only in DataFrame.corr is deprecated. In a future version, it will  
default to False. Select only valid columns or specify the value of numeric_only  
to silence this warning.  
    corr_matrix = data.corr()
```

Correlation Matrix





In [... #model 1

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.feature_selection import SelectKBest, f_classif
from keras.models import Sequential
from keras.layers import Dense
import matplotlib.pyplot as plt
import numpy as np

data = pd.read_csv('/content/Brain Tumor.csv')
data = data.replace([np.inf, -np.inf], np.nan)
data = data.fillna(data.mean())

X = data.drop(['Image', 'Class'], axis=1)
y = data['Class']

# Feature engineering
selector = SelectKBest(score_func=f_classif, k=10)
X_selected = selector.fit_transform(X, y)

selected_indices = selector.get_support(indices=True)

selected_columns = X.columns[selected_indices]

print("Selected Features:")
for column in selected_columns:
    print(column)

X_train, X_test, y_train, y_test = train_test_split(X_selected, y, test_size=0.2)

# Normalize
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

model = Sequential()
model.add(Dense(64, activation='relu', input_shape=(X_train.shape[1],)))
model.add(Dense(64, activation='relu'))
model.add(Dense(1, activation='sigmoid'))

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

history = model.fit(X_train, y_train, epochs=20, validation_split=0.2)

_, train_acc = model.evaluate(X_train, y_train, verbose=0)
_, test_acc = model.evaluate(X_test, y_test, verbose=0)

print('Train: %.3f, Test: %.3f' % (train_acc, test_acc))

plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'])
```

```
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')

plt.show()
```

<ipython-input-31-a00c492d20cc>:12: FutureWarning: The default value of numeric_only in DataFrame.mean is deprecated. In a future version, it will default to False. In addition, specifying 'numeric_only=None' is deprecated. Select only valid columns or specify the value of numeric_only to silence this warning.

```
data = data.fillna(data.mean())
```

Selected Features:

Variance

Standard Deviation

Entropy

Skewness

Kurtosis

Contrast

Energy

ASM

Homogeneity

Dissimilarity

Epoch 1/20

76/76 [=====] - 1s 7ms/step - loss: 0.2654 - accuracy: 0.9418 - val_loss: 0.0817 - val_accuracy: 0.9784

Epoch 2/20

76/76 [=====] - 0s 5ms/step - loss: 0.0771 - accuracy: 0.9767 - val_loss: 0.0515 - val_accuracy: 0.9850

Epoch 3/20

76/76 [=====] - 0s 5ms/step - loss: 0.0644 - accuracy: 0.9805 - val_loss: 0.0460 - val_accuracy: 0.9884

Epoch 4/20

76/76 [=====] - 0s 5ms/step - loss: 0.0595 - accuracy: 0.9801 - val_loss: 0.0427 - val_accuracy: 0.9884

Epoch 5/20

76/76 [=====] - 0s 5ms/step - loss: 0.0547 - accuracy: 0.9817 - val_loss: 0.0399 - val_accuracy: 0.9900

Epoch 6/20

76/76 [=====] - 0s 6ms/step - loss: 0.0510 - accuracy: 0.9846 - val_loss: 0.0397 - val_accuracy: 0.9917

Epoch 7/20

76/76 [=====] - 0s 5ms/step - loss: 0.0475 - accuracy: 0.9855 - val_loss: 0.0388 - val_accuracy: 0.9900

Epoch 8/20

76/76 [=====] - 0s 5ms/step - loss: 0.0446 - accuracy: 0.9838 - val_loss: 0.0398 - val_accuracy: 0.9900

Epoch 9/20

76/76 [=====] - 0s 5ms/step - loss: 0.0434 - accuracy: 0.9842 - val_loss: 0.0368 - val_accuracy: 0.9917

Epoch 10/20

76/76 [=====] - 1s 17ms/step - loss: 0.0416 - accuracy: 0.9850 - val_loss: 0.0380 - val_accuracy: 0.9900

Epoch 11/20

76/76 [=====] - 0s 6ms/step - loss: 0.0408 - accuracy: 0.9846 - val_loss: 0.0359 - val_accuracy: 0.9884

Epoch 12/20

76/76 [=====] - 0s 3ms/step - loss: 0.0380 - accuracy: 0.9863 - val_loss: 0.0344 - val_accuracy: 0.9884

Epoch 13/20

76/76 [=====] - 0s 4ms/step - loss: 0.0381 - accuracy: 0.9859 - val_loss: 0.0375 - val_accuracy: 0.9900

Epoch 14/20

76/76 [=====] - 0s 3ms/step - loss: 0.0372 - accuracy: 0.9863 - val_loss: 0.0340 - val_accuracy: 0.9884

Epoch 15/20


76/76 [=====] - 0s 4ms/step - loss: 0.0356 - accuracy: 0.9880 - val_loss: 0.0317 - val_accuracy: 0.9917

Epoch 16/20

76/76 [=====] - 0s 4ms/step - loss: 0.0336 - accuracy: 0.9867 - val_loss: 0.0343 - val_accuracy: 0.9900

Epoch 17/20

76/76 [=====] - 0s 3ms/step - loss: 0.0318 - accuracy: 0.9884 - val_loss: 0.0318 - val_accuracy: 0.9917



```
In [... from sklearn.metrics import precision_score, recall_score, f1_score, roc_auc
```

```
y_pred = model.predict(X_test)
y_pred = [1 if p > 0.5 else 0 for p in y_pred]
```

```
#performance metrics
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred)
```

```
print('Precision: %.3f, Recall: %.3f, F1 Score: %.3f, ROC AUC: %.3f' % (prec
```

```
24/24 [=====] - 0s 1ms/step
```

```
Precision: 0.991, Recall: 0.973, F1 Score: 0.982, ROC AUC: 0.983
```

```
In [33]: import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix, roc_curve, roc_auc_score
```

```
y_pred = (model.predict(X_test) > 0.5).astype("int32")
y_pred_proba = model.predict(X_test)
```

```
plt.figure(figsize=(10,10))
sns.heatmap(data.corr(), annot=True, fmt=".2f")
plt.show()
```

```
plt.figure(figsize=(10,6))
plt.title('Distribution of Predictions')
sns.histplot(y_pred_proba, kde=True)
plt.show()
```

```
cf_matrix = confusion_matrix(y_test, y_pred)
sns.heatmap(cf_matrix, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```


```
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)
auc = roc_auc_score(y_test, y_pred_proba)
plt.plot(fpr, tpr, label="AUC="+str(auc))
plt.legend(loc=4)
plt.title('Receiver Operating Characteristic Curve')
plt.show()
```

```
24/24 [=====] - 0s 1ms/step
```


```
24/24 [=====] - 0s 2ms/step
```

<ipython-input-33-39b6281ba315>:10: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.


```
sns.heatmap(data.corr(), annot=True, fmt=".2f")
```




Distribution of Predictions



Confusion Matrix



Receiver Operating Characteristic Curve



In [3... #Image Classificaition Model #Model2

```
import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D
from tensorflow.keras.models import Model
import os

from google.colab import drive
drive.mount('/content/drive')

dataset_directory = '/content/drive/MyDrive/brain_tumor_dataset'

image_size = (330, 330)

datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=30,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    validation_split=0.2
)

data_generator = datagen.flow_from_directory(
    dataset_directory,
    target_size=image_size,
    batch_size=32,
    class_mode='binary',
    subset='training'
)

validation_generator = datagen.flow_from_directory(
    dataset_directory,
    target_size=image_size,
    batch_size=32,
    class_mode='binary',
    subset='validation'
)

base_model = MobileNetV2(
    include_top=False,
    weights='imagenet',
    input_shape=(image_size[0], image_size[1], 3)
)

base_model.trainable = False

x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(64, activation='relu')(x)
predictions = Dense(1, activation='sigmoid')(x)

model = Model(inputs=base_model.input, outputs=predictions)
```

```
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

history = model.fit(
    data_generator,
    steps_per_epoch=data_generator.samples // data_generator.batch_size,
    validation_data=validation_generator,
    validation_steps=validation_generator.samples // validation_generator.batch_size,
    epochs=10
)
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

Found 203 images belonging to 2 classes.

Found 50 images belonging to 2 classes.

WARNING:tensorflow:`input_shape` is undefined or non-square, or `rows` is not in [96, 128, 160, 192, 224]. Weights for input shape (224, 224) will be loaded as the default.

Epoch 1/10

6/6 [=====] - 47s 7s/step - loss: 0.7217 - accuracy: 0.5833 - val_loss: 0.6014 - val_accuracy: 0.7188

Epoch 2/10

6/6 [=====] - 33s 5s/step - loss: 0.5737 - accuracy: 0.7310 - val_loss: 0.5025 - val_accuracy: 0.6562

Epoch 3/10

6/6 [=====] - 31s 5s/step - loss: 0.5325 - accuracy: 0.7310 - val_loss: 0.4072 - val_accuracy: 0.8438

Epoch 4/10

6/6 [=====] - 33s 6s/step - loss: 0.4869 - accuracy: 0.7953 - val_loss: 0.3066 - val_accuracy: 0.8750

Epoch 5/10

6/6 [=====] - 36s 6s/step - loss: 0.4452 - accuracy: 0.7895 - val_loss: 0.2711 - val_accuracy: 0.9062

Epoch 6/10

6/6 [=====] - 37s 7s/step - loss: 0.4377 - accuracy: 0.8070 - val_loss: 0.2085 - val_accuracy: 1.0000

Epoch 7/10

6/6 [=====] - 38s 6s/step - loss: 0.3843 - accuracy: 0.8480 - val_loss: 0.2709 - val_accuracy: 0.9375

Epoch 8/10

6/6 [=====] - 30s 5s/step - loss: 0.3414 - accuracy: 0.8655 - val_loss: 0.2850 - val_accuracy: 0.9062

Epoch 9/10

6/6 [=====] - 32s 5s/step - loss: 0.3518 - accuracy: 0.8538 - val_loss: 0.1738 - val_accuracy: 0.9375

Epoch 10/10

6/6 [=====] - 32s 5s/step - loss: 0.3535 - accuracy: 0.8772 - val_loss: 0.2736 - val_accuracy: 0.8750

In [35... #testing accuracy on validation set

```
import pandas as pd
```

```
loss, accuracy = model.evaluate(validation_generator)
print(f'Loss: {loss}, Accuracy: {accuracy}')
```

```
predictions = model.predict(validation_generator)
```

```
binary_predictions = [1 if p > 0.5 else 0 for p in predictions.flatten()]
```

```
df_predictions = pd.DataFrame(binary_predictions, columns=['Predictions'])
```

```
print(df_predictions)
```


2/2 [=====] - 8s 3s/step - loss: 0.2583 - accuracy: 0.9200

Loss: 0.2582961618900299, Accuracy: 0.9200000166893005

2/2 [=====] - 11s 1s/step

Predictions

0	1
1	0
2	1
3	0
4	0
5	1
6	1
7	0
8	0
9	0
10	0
11	0
12	0
13	0
14	1
15	1
16	1
17	0
18	0
19	0
20	1
21	1
22	0
23	0
24	1
25	0
26	1
27	1
28	1
29	1
30	0
31	1
32	1
33	1
34	1
35	1
36	0
37	1
38	0
39	1
40	1
41	1
42	1
43	1
44	0
45	0
46	1
47	1
48	0
49	1

In [3... #prediction on unseen data

```
import os
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import pandas as pd

unseen_dataset_directory = '/content/drive/MyDrive/unseen3'

unseen_datagen = ImageDataGenerator(rescale=1./255)

unseen_data_generator = unseen_datagen.flow_from_directory(
    unseen_dataset_directory,
    target_size=image_size,
    batch_size=32,
    class_mode=None,
    shuffle=False
)
file_names = unseen_data_generator.filenames

predictions = model.predict(unseen_data_generator)

binary_predictions = [1 if p > 0.5 else 0 for p in predictions.flatten()]

df_predictions = pd.DataFrame({'Image': file_names, 'Prediction': binary_pr

#print(df_predictions)
```

Found 205 images belonging to 2 classes.

7/7 [=====] - 25s 3s/step

```
In [... df_predictions['TrueLabel'] = df_predictions['Image'].apply(lambda x: 1 if '

correct_predictions = (df_predictions['TrueLabel'] == df_predictions['Predic

accuracy = correct_predictions / len(df_predictions)

print(f'Accuracy: {accuracy * 100}%')

Accuracy: 82.92682926829268%
```

```
In [3... import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve, auc
```

```
plt.figure(figsize=(10, 6))
```

```
df_predictions['Prediction'].value_counts().plot(kind='bar', color=['salmon
plt.title('Predictions')
plt.xlabel('Class')
plt.ylabel('Count')
plt.xticks([0, 1], ['No Tumor', 'Glioma Tumor'], rotation=0)
plt.show()
```

